# Chapter ONE Probably Approximately Correct (PAC)

Siheng Zhang

zhangsiheng@cvte.com

September 3, 2020

*The notes is mainly based on the following book*

- Understanding Machine Learning: From Theory to Algorithms, Shai Shalev-Shwartz and Shai Ben-David, 2014 [1]

- pattern recognition and machine learning, Christopher M. Bishop, 2006 [2]

- Probabilistic Graphical Models: Principles and Techniques, Daphne Koller and Nir Friedman, 2009 [3]

- Graphical Models, Exponential Families, and Variational Inference, Martin J. Wainwright and Michael I. Jordan, 2008 [4]

*Corresponding to Chapter 2-5 in UML.*
This part mainly answers the following questions:

- What can we know about the generalization error?

- How does the hypothesis set (in application, the choice of classifier/regressor or so on) reflect our prior knowledge, or, inductive bias?

---

[1] https://www.cs.huji.ac.il/s̄hais/UnderstandingMachineLearning/understanding-machine-learning-theory-algorithms.pdf

[2] http://users.isr.ist.utl.pt/w̄urmd/Livros/school/Bishop%20-%20Pattern%20Recognition%20And%20Machine%20Learning%20-%20Springer%20%202006.pdf

[3] https://mitpress.mit.edu/books/probabilistic-graphical-models

[4] https://people.eecs.berkeley.edu/w̄ainwrig/Papers/WaiJor08_FTML.pdf

# Contents

# 1 Formulation

## 1.1 The learner's input, output, and evaluation

- **input**:

  - Domain Set: instance $x \in \mathcal{X}$.
  - Label Set: label $y \in \mathcal{Y}$. Currently, just consider the binary classification task.
  - Training data: $S = ((x_1, y_1), \cdots, (x_m, y_m))$ is a finite sequence.

- **output**: hypothesis (or classifier, regressor) $h : \mathcal{X} \to \mathcal{Y}$.

- **data generation model**: Assume that the instances are generated by some probability distribution $\mathcal{D}$, and there is some 'correct' labeling function (currently): $f : \mathcal{X} \to \mathcal{Y}$.

  The i.i.d. assumption: the training samples are independently and identically distributed.

  *remark1: The learner is blind to the data generation model.*

  *remark2: Usually called 'training set', but must be 'training sequence', because the same sample may repeat, and some training algorithms is order-sensitive.*

  *remark3: Strictly speaking, the distribution $\mathcal{D}$ is defined over $\mathcal{X} \times \mathcal{Y}$.*

- **Generalization error**: *a.k.a*, true error/risk.

$$L_{\mathcal{D},f}(h) \overset{def}{=} \underset{x \sim \mathcal{D}}{\mathbb{P}}[h(x) \neq f(x)] \overset{def}{=} \mathcal{D}(x : h(x) \neq f(x)) \tag{1}$$

# 2 From ERM to PAC

## 2.1 ERM (Empirical Risk Minimization) may lead to overfitting

Since the generalization error is intractable, turn to minimize the **empirical risk**:

$$L_S(h) \overset{def}{=} \frac{|\{(x_i, y_i) \in S : h(x_i) \neq y_i\}|}{m} \tag{2}$$

Consider a 'lazy' learner $h$, which predict $y = y_i$ iff. $x = x_i$, and 0 otherwise. It has $1/2$ probability to fail for unseen instances, i.e., $L_{\mathcal{D},f}(h) = 1/2$, while $L_S(h) = 0$. Hence, it is an excellent learner on the training set, but a poor learner in the universe case. This phenomenon is called 'overfitting'. The lesson behind this learner is: without restriction on the hypothesis set, ERM can lead to overfitting.

## 2.2 ERM with restricted hypothesis set (inductive bias)

Instead of $h_S \in \arg\min L_S(h)$, ERM with restricted hypothesis set return the following hypothesis:

$$h_S \in \arg\min_{h \in \mathcal{H}} L_S(h) \tag{3}$$

Start from an ideal case, in which the **realizability assumption** holds, i.e., there exists $h^* \in \mathcal{H}$, such that $L_{\mathcal{D},f}(h^*) = 0$.

It implies that $L_S(h^*) = 0$, $L_S(h_S) = 0$. However, we are interested in $L_{\mathcal{D},f}(h_S)$.

## 2.3 PAC (Probably Approximately Correct) learnability

**Definition**: Training on $m \geq m_{\mathcal{H}}(\epsilon, \delta)$ samples, there exists an algorithm to be able to achieve **accuracy** at least $1 - \epsilon$ with **confidence** at least $1 - \delta$.

**Theorem 1** *Finite hypothesis classes are PAC learnable, and the sample complexity is: $m_{\mathcal{H}}(\epsilon, \delta) = \frac{\log(|\mathcal{H}|/\delta)}{\epsilon}$.*

**Proof** *Let $\mathcal{H}_B$ be the set of 'bad' hypothesis, that is, $\mathcal{H}_B \subset \mathcal{H}$, and $\forall h \in \mathcal{H}_B, L_{\mathcal{D},f}(h) > \epsilon$. Let $M$ be the set of 'misleading' samples, that is $M = \{S : \exists h \in \mathcal{H}_B, L_S(h) = 0\}$. Note that,*

$$M = \bigcup_{h \in \mathcal{H}_B} \{S : L_S(h) = 0\} \tag{4}$$

*The goal is to bound the probability of the event $L_{\mathcal{D},f}(h_S) > \epsilon$,*

$$
\begin{aligned}
\mathcal{D}^m(\{S : L_{\mathcal{D},f}(h_S) > \epsilon\}) &\leq \mathcal{D}^m(M) \\
&= \mathcal{D}^m(\bigcup_{h \in \mathcal{H}_B} \{S : L_S(h) = 0\}) = \sum_{h \in \mathcal{H}_B} \prod_{i=1}^{m} \mathcal{D}(\{x_i : f(x_i) = h(x_i)\}) \\
&\overset{i.i.d.}{=} \sum_{h \in \mathcal{H}_B} (1 - L_{\mathcal{D},f}(h))^m \leq \sum_{h \in \mathcal{H}_B} (1 - \epsilon)^m \leq \sum_{h \in \mathcal{H}_B} \exp(-\epsilon m) \\
&\leq |\mathcal{H}| \exp(-\epsilon m)
\end{aligned}
\tag{5}
$$

*Let $|\mathcal{H}| \exp(-\epsilon m) \leq \delta$, we can solve that $m \geq \log(|\mathcal{H}|/\delta)/\epsilon$.*

## 2.4 No-Free-Lunch

**Theorem 2** *Let $A$ be any learning algorithm for the task of binary classification with respect to the 0-1 loss over a domain $\mathcal{X}$. Let $m$ be any number smaller than $\mathcal{X}/2$, representing a training set size. Then, there exists a distribution $\mathcal{D}$ over $X \times \{0,1\}$ such that:*

- *There exists a function $f : \mathcal{X} \to \{0,1\}$ with $L_{\mathcal{D}}(f) = 0$.*

- *With probability of at least 1/7 over the choice of $S \sim \mathcal{D}^m$ we have that $L_{\mathcal{D}}(A(S)) \geq 1/8$.*

NFL theorem tells the neccessity of inductive bias. Philosophically, if someone can explain every phenomenon, his explanations are worthless.

## 2.5 Agnostic PAC

### 2.5.1 Beyond realizability assumption

In practical, the 'true' labelling function may not exist, and the labels may not be fully determined by the features on hand. Then Agnostic PAC learnability is defined as: training on $m \geq m_{\mathcal{H}}(\epsilon, \delta)$ samples, there exists an algorithm with **confidence** at least $1 - \delta$ to achieve that:

$$L_{\mathcal{D}}(h) \leq \min_{h' \in \mathcal{H}} L_{\mathcal{D}}(h') + \epsilon \tag{6}$$

in which $L_{\mathcal{D}}(h) \overset{def}{=} \mathbb{P}_{(x,y) \sim \mathcal{D}}[h(x) \neq y] \overset{def}{=} \mathcal{D}(\{x : h(x) \neq y\})$.

### 2.5.2 Beyond binary classification

Agnostic PAC learnability remains the same with:

$$\mathcal{D}(h) = \mathbb{E}_{x \sim \mathcal{D}}[l(h, z)] \tag{7}$$

in which $l(\cdot)$ is 0-1 loss for multiclass classification and square loss for regression.

### 2.5.3 Sample complexity under Agn-PAC: via uniform convergence

# 3 Error decomposition

# 4 Summary

Now that, we have come to some important conclusions under the PAC learning framework:

1. No universal learner;

2. Inductive bias is neccessary to avoid overfitting;

3. Sample complexity is function about hypothesis set, confidence level and error, interestingly, it is nothing to do with the dimension of feature space;

4. Inductive bias controls the balance of approximation error and estimation error.

We have reached the fundamental question in learning theory: **Over which hypothesis classes, ERM learning will not result in overfitting (or, PAC learnable)?** Currently, we just confirm the PAC learnability for finite classes. In the next chapter, the most important part in learning theory, VC-dimension, will gives a more precise answer.

# 5    Excercises and solutions