Chapter 4 Linear Model

Siheng Zhang
zhangsiheng@cvte.com

December 15, 2020

This part corresponds to **Chapter 1,3,4 of PRML, Chapter of UML**, and mainly answers the following questions:

-

Contents

1 Linear classification

In the last chapter, we stops at the linear classification of binary classification task,

$$y = h(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + w_0 \tag{1}$$

in which $\mathbf{w}$ is weight vector, and w_0 is bias. The input vector is assigned to class C_1 iff. $h(\mathbf{x}) \geq 0$ and to class C_2 otherwise.

Consider two points $\mathbf{x}_1, \mathbf{x}_2$ on the decision boundary, i.e., $\mathbf{w}^\top(\mathbf{x}_1 - \mathbf{x}_2) = 0$, hence $\mathbf{w}$ is orthogonal to the decision boundary. And the distance from the origin to the decision boundary is

$$\frac{\mathbf{w}^\top \mathbf{x}}{\|\mathbf{w}\|} = \frac{-w_0}{\|\mathbf{w}\|} \tag{2}$$

It is usually convenient to use a more compact notation in which we introduce an additional input value $x_0 = 1$ and then define $\tilde{\mathbf{w}} = (w0, \mathbf{w})$ and $\tilde{\mathbf{x}} = (x0, \mathbf{x})$ so that $y = f(\mathbf{x}) = \tilde{\mathbf{w}}^\top \tilde{\mathbf{x}}$.

1.1 Extend to multiple classes

1.2 Fisher's linear discriminant

2 Linear regression

$$\min_{\mathbf{w}} L_S(h) = \sum_{i=1}^{m} (h(\mathbf{x}_i) - y_i)^2 = \sum_{i=1}^{m} (\mathbf{w}\mathbf{x}_i - y_i)^2 \tag{3}$$

Suppose the fitting error $\epsilon_i = y_i - \mathbf{w}\mathbf{x}_i$ is Gaussian noise, i.e., $\epsilon_i \sim \mathcal{N}(0, \beta)$. Then the log likelihood function of the training sequence is:

$$\log \mathcal{L} = -\frac{m}{2} \log 2\pi\beta - \sum_{i=1}^{m} \frac{(y_i - \mathbf{w}\mathbf{x}_i)^2}{2\beta} \tag{4}$$

Obviously, MLE is equivalent to linear regression.

remark1: Since linear regression is not a binary prediction task, we cannot analyse its sample complexity using the VC-dimension. One possible analysis of the sample complexity of linear regression is by relying on the "discretization trick". However, to apply the sample complexity bounds from Chapter 2 we also need that the loss function will be bounded.

over-fitting

2.1 Ridge regression

Ridge regression addresses on over-fitting by penalizing the l_2-norm of weight vector $\mathbf{w}$,

$$\min_{\mathbf{w}} \sum_{i=1}^{N} (\mathbf{w}\mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|_2^2$$

If we assume a Gaussian prior for the weight vector, $\mathbf{w} \sim \mathcal{N}(0, \alpha^{-1}\mathbf{I})$, then the posterior of the training sequence is:

$$p(\mathbf{w}|S) \propto p(\mathbf{w})p(S|\mathbf{w}) \propto \exp\left(-\frac{\alpha}{2}\mathbf{w}^\top \mathbf{w}\right) \cdot \prod_{i=1}^{N} \exp\left(-\frac{(y_i - \mathbf{w}\mathbf{x}_i)^2}{2\beta}\right) \tag{5}$$

Maximizing the log posterior function is equivalent to the ridge regression.

2.2 Lasso

Lasso addresses on over-fitting by penalizing the l_1-norm of weight vector $\mathbf{w}$,

$$\min_{\mathbf{w}} \sum_{i=1}^{N} (\mathbf{w}\mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|_1$$

If we assume a Laplace prior for the weight vector, $p(\mathbf{w}) = \frac{1}{2\alpha} \exp\left(-\frac{\|\mathbf{w}\|_1}{\alpha}\right)$, then the posterior of the training sequence is:

$$p(\mathbf{w}|S) \propto p(\mathbf{w})p(S|\mathbf{w}) \propto \exp\left(-\frac{\|\mathbf{w}\|_1}{\alpha}\right) \cdot \prod_{i=1}^{N} \exp\left(-\frac{(y_i - \mathbf{w}\mathbf{x}_i)^2}{2\beta}\right) \tag{6}$$

Maximizing the log posterior function is equivalent to the Lasso model.

3 Generalized linear model